

Supervised Learning with Quantum Computers

Isidora Araya Day
Radoica Draškić
Brennan Undseth

Delft University of Technology, The Netherlands

April 23, 2020

Overview


- 1 Quick introduction to Machine Learning
- 2 Quantum speedups?
- 3 Variational quantum circuits
- 4 Implementation on Quantum Hardware

Table of Contents

- 1 Quick introduction to Machine Learning
- 2 Quantum speedups?
- 3 Variational quantum circuits
- 4 Implementation on Quantum Hardware





What is Machine Learning?

- Supervised learning
 - Input labeled data
 - Train machine
 - Label new data
- Unsupervised learning
 - Analysis of unlabeled data
- Reinforcement learning



What is Machine Learning?

- Supervised learning
 - Input labeled data
 - Train machine
 - Label new data
- Unsupervised learning
 - Analysis of unlabeled data
- Reinforcement learning



Supervised learning in breast cancer diagnosis

Discrimination between benign and malignant samples after fine needle biopsies ¹.





Figure: left:malignant, right: benign.

10 Parameters: radius, perimeter, area, compactness, smoothness, concavity, concave points, fractal dimension, texture of tumours.

¹M. Sewak, P. Vaidya, C. Chan, and Zhong-Hui Duan (2007). "SVM Approach to Breast Cancer Classification". In: pp. 32–37


Support vector machine (SVM) ²

Labeled data:

$$\{(\vec{x}_j, y_j) : \vec{x}_j \in \mathbb{R}^N, y_j = \pm 1\}_{j=1, \dots, M} \quad (1)$$

Purpose: to find a dividing hyperplane

$$y_j(\vec{w} \cdot \vec{x}_j + b) \geq 1 \quad (2)$$



²L. Saitta (1995). "Support-Vector Networks". In: 297, pp. 273–297

Table of Contents

- 1 Quick introduction to Machine Learning
- 2 Quantum speedups?
- 3 Variational quantum circuits
- 4 Implementation on Quantum Hardware

Quantum machine learning

Motivation: Internal product and inverse matrix calculation ³.

In theory: HHL algorithm to solve linear systems of eq ⁴.

Bottleneck: Encoding.

In practice: Hybrid algorithms

A feature map that is **hard to estimate classically** is an important part of creating a quantum advantage

³P. Rebentrost, M. Mohseni, and S. Lloyd (2014). "Quantum Support Vector Machine for Big Data Classification". In: 130503.September, pp. 1–5. DOI: 10.1103/PhysRevLett.113.130503

⁴A. W. Harrow, A. Hassidim, and S. Lloyd (2009). "Quantum Algorithm for Linear Systems of Equations". In: 150502.October, pp. 1–4. DOI: 10.1103/PhysRevLett.103.150502

IBM - Hybrid approach ⁵

2 different methods:

- Quantum Kernel Estimator
- Quantum Variational Classifier

⁵V. Havlíček, A. D. Córcoles, K. Temme, A. W. Harrow, A. Kandala, J. M. Chow, and J. M. Gambetta (2019). "Supervised learning with quantum-enhanced feature spaces". In: *Nature* 567.7747, 209–212. ISSN: 1476-4687. DOI: 10.1038/s41586-019-0980-2

IBM - Hybrid approach ⁵

2 different methods:

- Quantum Kernel Estimator
- Quantum Variational Classifier




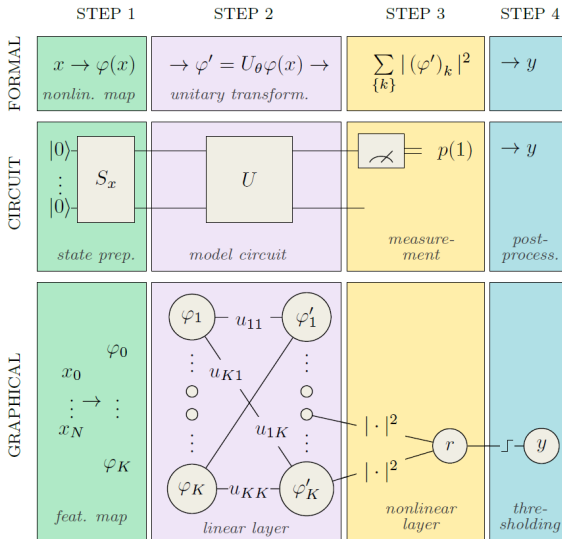
Figure: (a) 5 Qubit transmon circuit. (b) Results QKE Qiskit simulation.

⁵V. Havlíček, A. D. Córcoles, K. Temme, A. W. Harrow, A. Kandala, J. M. Chow, and J. M. Gambetta (2019). "Supervised learning with quantum-enhanced feature spaces". In: *Nature* 567.7747, 209–212. ISSN: 1476-4687. DOI: 10.1038/s41586-019-0980-2

Table of Contents

- ① Quick introduction to Machine Learning
- ② Quantum speedups?
- ③ Variational quantum circuits
- ④ Implementation on Quantum Hardware

Workflow⁶



⁶M. Schuld, A. Bocharov, K. M. Svore, and N. Wiebe (2020). "Circuit-centric quantum classifiers". In: *Phys. Rev. A* 101 (3), p. 032308. DOI: 10.1103/PhysRevA.101.032308

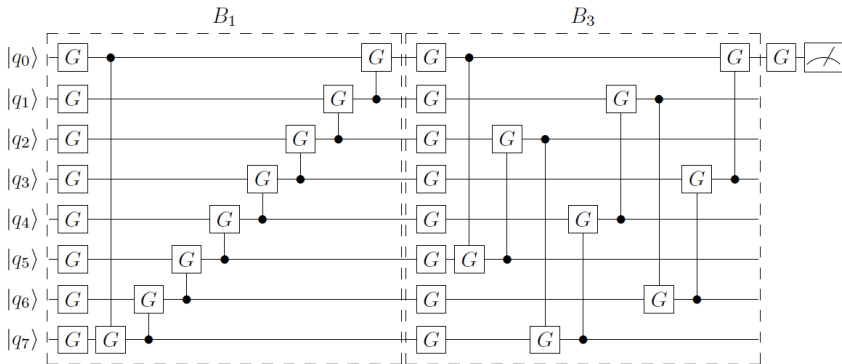
State preparation

$$\begin{pmatrix} x_1 \\ x_2 \\ \vdots \\ x_N \end{pmatrix} \rightarrow \frac{1}{\sqrt{\sum_j x_j^2 + \sum_k |c_k|^2}} \begin{pmatrix} x_1 \\ \vdots \\ x_N \\ c_1 \\ \vdots \\ c_D \end{pmatrix} \equiv |\varphi(x)\rangle$$

$$|\varphi(x)\rangle \rightarrow \underbrace{|\varphi(x)\rangle \otimes \cdots \otimes |\varphi(x)\rangle}_{d \text{ times}}$$

Tensorial feature map will hopefully introduce nonlinearities that may facilitate the classification procedure.

Variational circuit architecture




17 single qubit gates $G(\alpha, \beta, \gamma)$, 16 two qubit gates $CG(\alpha, \beta, \gamma)$.

Total number of parameters $33 \times 3 + 1 = 100$.

Training the model

- Prediction $\pi(x; \theta, b) = \langle \varphi(x) | U^\dagger(\theta) (Z \otimes \cdots \otimes I) U(\theta) | \varphi(x) \rangle + b$
- Cost function $\mathcal{C}(\theta, b) = \frac{1}{2M} \sum_{m=1}^M |\pi(x^m; \theta, b) - y^m|^2$
- Gradient descent $\theta_i^{(t)} = \theta_i^{(t-1)} - \eta \partial \mathcal{C} / \partial \theta_i$, $b^{(t)} = b^{(t-1)} - \eta \partial \mathcal{C} / \partial b$




 PENNYLANCE

 Qiskit

Linearly separable data


- 2 qubits
- padding with 2 uninformative features
- 97% accuracy



$$(x \ y)^T \rightarrow (x \ y \ 1 \ 1)^T$$


Concentric circles

- 4 qubits
- padding with 2 uninformative features
- $d = 2$
- 99% accuracy




$$\begin{aligned} \begin{pmatrix} x & y \end{pmatrix}^T &\rightarrow \begin{pmatrix} x & y & 1 & 1 \end{pmatrix}^T \otimes \begin{pmatrix} x & y & 1 & 1 \end{pmatrix}^T \\ &= \begin{pmatrix} x^2 & xy & x & x & xy & y^2 & y & y & \dots \end{pmatrix}^T \end{aligned}$$

Moons



2 qubits
 $d = 1$
80% accuracy

$$(x \ y)^T \rightarrow (x \ y \ 1 \ 1)^T$$




4 qubits
 $d = 2$
96% accuracy

$$(x \ y)^T \rightarrow (x \ y \ 1 \ 1)^{\otimes 2, T}$$

Table of Contents

- 1 Quick introduction to Machine Learning
- 2 Quantum speedups?
- 3 Variational quantum circuits
- 4 Implementation on Quantum Hardware

Exclusive-OR Data



We use the training data to optimize the variational parameters $\vec{\theta}$ with respect to a **cost** function.


We use the test data to quantify the **accuracy** of our model.

State Preparation S_x

- Input in Cartesian coordinates:
(x,y)
- Encode state in ket:


$$\frac{1}{\chi} \begin{pmatrix} x \\ y \end{pmatrix} \otimes \begin{pmatrix} x \\ y \end{pmatrix} = \frac{1}{\chi} \begin{pmatrix} x^2 \\ xy \\ yx \\ y^2 \end{pmatrix}$$

- Note that states in quadrants II and IV pick up relative phase




Variational Circuit $U(\vec{\theta})$

Since IBM Quantum Experience does not support arbitrary controlled unitaries...



We expand to the following circuit:



- **Input** datum x^m with label y^m
- **Prediction** $\pi^m = \langle 0 | S_x^\dagger U^\dagger(\vec{\theta}) Z_1 U(\vec{\theta}) S_x | 0 \rangle + b$
- **Optimize** $\vec{\theta}, b$ to minimize the **cost** function:
$$C = \frac{1}{2M} \sum_i^M (\pi^m - y^m)^2$$
- State preparation and optimization is handled easily with PennyLane



P E N N Y L A N E

```
@qml.qnode(dev)
def circuit(weights, x=None):
    statepreparation(x)

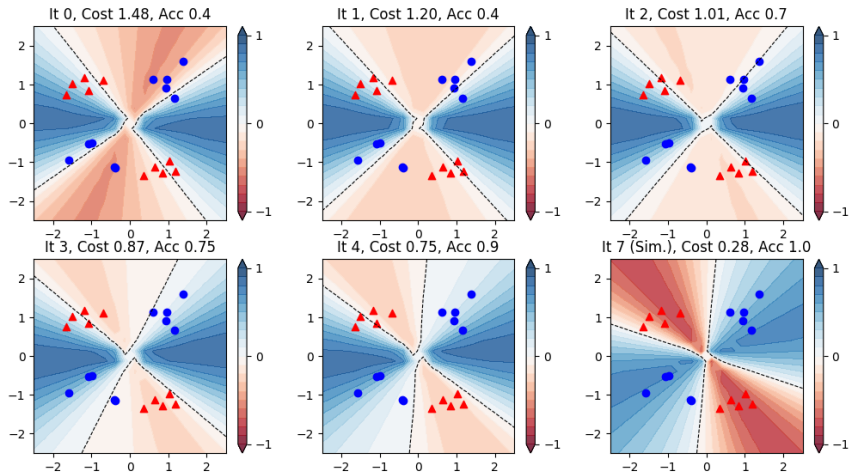
    qml.Rot(weights[0, 0], weights[0, 1], weights[0, 2], wires=0)
    qml.Rot(weights[1, 0], weights[1, 1], weights[1, 2], wires=1)
    qml.CNOT(wires=[1,0])
    qml.Rot(weights[2, 0], weights[2, 1], weights[2, 2], wires=0)
    qml.CNOT(wires=[1,0])
    qml.Rot(weights[3, 0], weights[3, 1], weights[3, 2], wires=0)

    return qml.expval(qml.PauliZ(0))

def model(var, x=None):
    weights, bias = var
    return circuit(weights, x=x) + bias

def loss(labels, predictions):
    loss = 0
    for l, p in zip(labels, predictions):
        loss = loss + (l - p) ** 2
    loss = loss / len(labels)
    return loss
```


Results with IBM Quantum Experience



IBM Q[™]

- IBM queue takes anywhere between seconds to hours
- Single iteration: 20 data points with 12 parameters requires 480 unique circuits
- Each iteration takes well over 1 day with 5 minute queue time
- Python script run on an Amazon cloud computer for a week


Conclusion

- Quantum computers can perform supervised machine learning tasks with existing quantum hardware, but dedicated computers will be required for meaningful tasks
- State encoding and optimization are non-trivial design choices and greatly influence classifier performance
- Many approaches and opportunities for quantum speedup, we have presented just one way

Thank you!
Questions?

Calculating the gradient

Shift rule



If f is some linear function of the expectation value $\langle \hat{B} \rangle$, and $G(\alpha, \beta, \gamma)$

$$\frac{\partial f(G(\mu))}{\partial \mu} = \frac{1}{2} [f(G(\mu + \pi/2)) - f(G(\mu - \pi/2))], \mu \in \{\alpha, \beta, \gamma\}.$$

Same result holds for $CG(\alpha, \beta, \gamma)$.

Kernel matrix:

$$K(\vec{x}, \vec{z}) = |\langle \Phi(\vec{x}) | \Phi(\vec{z}) \rangle|^2$$

$$U_{\Phi(x)} = \exp \left(i \sum_{S \subseteq [2]} \phi_S(x) \Pi_{i \in S} Z_i \right)$$

$$\phi_{\{i\}} = x_i \text{ and } \phi_{\{1,2\}} = (\pi - x_1)(\pi - x_2)$$

